Product Metrics, MVP and Launch: A Practical Guide

The North Star metric

A North Star metric is the single number that best captures the value your product delivers to users. Good ones measure customer value, not vanity, such as weekly active teams for a collaboration tool rather than total signups. The North Star aligns the whole team, because every project can be asked one question: does this move the North Star, and if not, why are we doing it.

Pirate metrics (AARRR)

AARRR breaks the user journey into five stages: Acquisition, Activation, Retention, Referral, and Revenue. Acquisition is how people find you. Activation is the first moment they feel value. Retention is whether they come back. Referral is whether they tell others. Revenue is whether they pay. Most teams over-invest in acquisition and under-invest in activation and retention, which is where growth actually compounds.

Leading versus lagging indicators

A lagging indicator, such as monthly revenue, tells you what already happened but is slow to move. A leading indicator, such as the number of new users who complete setup in their first session, predicts the lagging number and moves quickly. After a launch, watch leading indicators first, because they tell you within days whether the lagging results will follow.

Defining a minimum viable product

A minimum viable product is the smallest thing you can build to learn whether the idea works. The emphasis is on learning, not on being small for its own sake. A good MVP tests the riskiest assumption directly. If the risk is 'will people pay', a landing page with a price can be a better MVP than a half-built app.

Go-to-market and launch checklist

A launch is not just shipping code. Before launch, agree the goal and the metric that defines success, prepare support and documentation, brief the sales and marketing teams, and decide the rollout plan. A staged rollout to a small percentage of users first lets you catch problems before everyone sees them. Always know in advance what number would make you roll the launch back.

A/B testing and experimentation

An A/B test shows one version to one group and another version to a comparable group, then compares a chosen metric. It is the cleanest way to know whether a change actually caused an improvement rather than coinciding with one. Decide the metric and the sample size before you start, and let the test run its full course, because stopping early the moment results look good is how teams fool themselves.

Measuring launch success

Judge a launch against the goal you set before it, not against a number you pick afterwards to look good. Combine a leading indicator for early signal, an activation measure to see if new users reach value, and a retention measure a few weeks later to see if the value lasts. A launch that spikes signups but shows no retention has not succeeded; it has only borrowed attention.

Why retention matters most

Retention is the truest test of product value, because it measures whether people come back on their own. Acquisition can be bought, but retention has to be earned. A product with strong retention grows even with modest acquisition, since existing users stay and refer others, while a product with weak retention leaks users as fast as it adds them, no matter how good the launch looked on day one.